

Neotoma paper

Simon Goring

09 August, 2014

neotoma: A Programmatic Interface to the Neotoma Paleocological Database

Abstract:

Paleoecological data is an integral part of ecological analysis. It provides an opportunity to study vegetation and climate interactions at time scales that cannot be observed through modern field studies, and allows us to observe changes in ecological processes that operate at centennial and millennial scales. Paleoecological inference also can allow us to understand ecological processes in the absence of widespread anthropogenic influence.

Here we describe the R package `neotoma`, to be used to obtain and manipulate paleoecological data from the Neotoma Paleocological Database (<http://www.neotomadb.org>). `neotoma` searches the Neotoma Database for datasets associated with location, taxa or dataset types using the database's Application Programming Interface. The package can return full datasets or metadata associated with sites and provides the ability to standardize taxonomies using one of several recognized standard taxonomies from the published literature.

To assist with the use of the package we provide examples of key functions using examples from the published literature, for both plant and mammal taxa.

Introduction

Paleoecological data is increasingly used to understand patterns of biogeographical, climatic and evolutionary change at multiple spatial and temporal scales. Paleocoinformatics ((Brewer *et al.* 2012; Uhen *et al.* 2013)) is increasingly providing tools to researchers across disciplines to access and use large datasets spanning thousands of years. These datasets may be used to provide better insight into patterns of biomass burning (Blarquez *et al.* 2013; Power *et al.*), regional vegetation change ((Blois *et al.* 2013; Blarquez *et al.* 2014)) or changes in physical processes over time ((Goring *et al.* 2012)). Critically, paleoecological data lags behind modern ecological cyber-infrastructure in regards to accessibility and extent. The increasing interest in uniting ecological and paleoecological data to understand modern ecological patterns and future responses (Dietl & Flessa 2011; Behrensmeyer & Miller 2012; Fritz *et al.* 2013) means that efforts to unite these two, seemingly independent data-streams will rely, in part, on more robust tools to access and synthesize paleoecological data.

The statistical software R ((R Core Team 2014)) is commonly used for analysis of paleoecological data and several packages in R exist for analysis (`analogue`: (Simpson 2007; Simpson & Oksanen 2013); `rioja`: (Juggins 2013), `Bchron`: (Parnell 2014), `paleofire`: (Group 2014)). Notwithstanding these packages, the use of extensive paleoecological resources within R has traditionally relied on *ad hoc* methods of obtaining and importing data. This has meant reliance on static online datasets such as the NOAA Paleoclimate repository or North American Modern Pollen Database, and on the distribution of individual datasets from author to analyst.

With an increasing push to provide paleoecological publications that include numerically reproducible results (e.g., (Goring *et al.* 2012, 2013; Gill *et al.* 2013)) it is important to provide tools that allow analysts to directly access dynamic datasets, and to provide tools to support reproducible workflows. The rOpenSci project has provided a number of tools that can directly interact with application programmatic interfaces (APIs) to access data from a number of databases including `rfishbase` (FishBase: (Boettiger *et al.* 2012)) and `taxize` (Encyclopedia of Life, iPlant/Taxosaurus and others: (Chamberlain & Szöcs 2013)) among others.

To illustrate use cases for the **neotoma** package we present examples drawn from the paleoecological literature to illustrate how **neotoma** provides the tools to perform research that is critical to understanding paleoecological change in the Pleistocene in an open and reproducible manner.

The **neotoma** package

Here we describe **neotoma**, an R package that acts as an interface between a large dynamic database (the Neotoma Paleoecological Database; <http://neotomadb.org>) and statistical tools in R. **neotoma** uses an API to send data requests to Neotoma, and then forms data objects that can interact with existing packages such as **analogue** and **rioja**. The **neotoma** package also includes tools to standardize pollen data across sample sites using a set of commonly accepted pollen taxa.

Examples

Macdonald and Cwynar ((1991)) used pollen percentage data for *Pinus* to map the northward migration of lodgepole pine (*Pinus contorta* var *latifolia*) following glaciation. In their study a cutoff of 15% *Pinus* pollen is associated with presence at pollen sample sites. Recent work by Strong and Hills ((Strong & Hills 2013)) has remapped the migration front using a lower pollen proportion (5%) and more sites. Here, we attempt to replicate the analysis as an example both of the strengths of the package and limitations of paleoinformatic approaches.

To begin we must define a spatial bounding box and a set of taxa of interest. Strong and Hills ((Strong & Hills 2013)) use a region approximately bounded by 54°N to the south and 65°N to the North, and from 110°W to 130°W. The function **get_site** is used to find all sites within a bounding box:

```
library("neotoma")
library("ggmap")
library("ggplot2")
library("reshape2")
library("plyr")
library("Bchron")
library("gridExtra")
all.sites <- get_site(loc = c(-140, 45, -110, 65))
```

The **get_sites** function returns a site data frame, with the columns **siteID**, **latitude**, **longitude**, **altitude**, **SiteName**, and **SiteDescription**. Each row represents a unique site.

We can see that this returns a total of ?? sites. Sites are effectively containers for datasets though. Generally it's better to search for datasets. When you search for a dataset you can limit the type of dataset, either by looking for specific taxa, or by describing the dataset type. Here we will look for all taxa beginning with *Pinus* in a pollen dataset. We use the ***** wildcard to indicate any and all taxa with *Pinus* in their name:

```
all.datasets <- get_dataset(loc = c(-140, 45, -110, 65), datasettype = "pollen",
  taxonname = "Pinus*")
```

A dataset is a larger data object. The dataset has site information, but it also has information about the specific dataset.

Here the API tells us we now have only ?? records of the original ?. Many of the samples are pollen surface samples, or vertebrate fauna, meaning pollen core data comprises less than half of the records. Regardless, we now know that there is pollen core data from ?? sites and we can plot those sites over our original ?.

```
map <- map_data("world")
ggplot(data = data.frame(map), aes(long, lat)) + geom_polygon(aes(group = group),
  color = "blue", alpha = 0.2) + geom_point(data = all.sites, aes(x = long,
  y = lat)) + geom_point(data = get_site(dataset = all.datasets), aes(x = long,
  y = lat), color = 2) + xlab("Longitude West") + ylab("Latitude North") +
  coord_map(projection = "albers", lat0 = 40, lat1 = 65, xlim = c(-140, -110),
  ylim = c(40, 70))
```

So we see that there are a number of sites in the interior of British Columbia that have no core pollen. For many of these cores pollen records exist. This is an obvious limitation of the use of large datasets. While many dataset have been entered into Neotoma, a large number have yet to make their way into the repository. An advantage of the API-based analysis however is that analysis using Neotoma can be updated continuously as new sites are added.

Let's get the data for each of the cores we have:

```
# This step may be time consuming when you run it, particularly if you have a
# slow internet connection.
all.downloads <- suppressMessages(get_download(dataset = all.datasets))
```

In most cases `get_download` will return a message for an individual core such as:

```
API call was successful. Returned record for Cottonwood Slough.
API call was successful. Returned record for Goring Woods.
```

The `download` object is a list with six components:

```
names(all.downloads[[1]])
```

The `metadata` component is again a list with a `dataset`, similar to the one returned by `get_dataset` plus a component, `pi.data`, that contains information about the principal investigator. The `sample.meta` component is where the core depth and age information is stored. The actual chronologies are stored in the `chronology` component. If a core has a single record this component has a length of one. Some cores have multiple chronologies and these are added to the list. The default chronology is always represented in `sample.meta`, and is always the first `chronology`. If you choose to build your own chronology using `Bacon` ((Blaauw *et al.* 2011)) or another method you can obtain the chronological controls for the core using the `get_chroncontrol` function and the chronology ID in either `sample.meta` or any one of the `chronology` objects. While the chronological controls used to build a chronology may vary across chronologies for a single site, the default model often contains the most accurate chronological control data.

The `taxon.list` component is a critical part of the `download` object. It lists the taxa found in the core, as well as any laboratory data, along with the units of measurement and taxonomic grouping. This is important information for determining which taxa make it into pollen percentages. The `counts` are the actual count or percentage data recorded for the core. The `lab.data` component contains information about any spike used to determine concentrations, sample quantities and, in some cases, charcoal counts.

We have ?? records in our analysis. Pollen taxonomy can vary substantially across cores often depending on researcher skill, or changing taxonomies for species, genera or families over time. This shifting taxonomy is often problematic to deal with. The `neotoma` package implements a taxonomic standardizer to attempt to standardize to one of four published taxonomies for the United States and Canada. While this function can be helpful in many cases it should also be used with care. The aggregation table is accessible using `data(pollen.equiv)` and the function to compile the data is called `compile_taxa`.

For our purposes we are really only interested in the percentage of *Pinus* in the core, so we can compile the taxa to the most straightforward taxonomy, 'P25' from Gavin *et al.* ((Gavin *et al.* 2003)). The first

record downloaded is Andy Lake, published by Szeicz ((Szeicz *et al.* 1995)). We can see in the `download` the `taxon.table` has 5 columns:

```
kable(head(all.downloads[[1]]$taxon.list))
```

Once we apply the `compile_taxa` function to the dataset using the ‘P25’ compiler:

```
compiled.cores <- lapply(all.downloads, compile_taxa, list.name = "P25")
```

we can see that the `taxon.table` now has an extra column (note that we’ve removed several columns to improve readability here).

```
kable(head(compiled.cores[[1]]$taxon.list[, c(1, 5, 6)]))
```

`compile_taxa` returns an object that looks exactly like the `download` object passed to it, however, the `taxon.list` data frame gains a column named `compressed` that links the original taxonomy to the revised taxonomy. This acts as an important check for researchers who choose to use this package for large-scale analysis. Here we see that taxa such as *Potentilla* is lumped into `Other`, along with spores and other taxa. The `compile_taxa` function can also accept user-defined tables for aggregation if the provided compilations are not acceptable.

In this case the counts look reasonable, and the synonymy appears to have been applied correctly (although we’re really only interested in *Pinus*). We now transform our `counts` into percentages to standardize across cores. We can see what a single core looks like:

```
# Get the percentage data for the first core:
core.pct <- as.data.frame(compiled.cores[[1]]$counts/rowSums(compiled.cores[[1]]$counts)) *
  100
core.pct$depth <- compiled.cores[[1]]$sample.meta$depths
core.pct$age <- compiled.cores[[1]]$sample.meta$Age
# Eliminate taxa with no samples greater than 5%.
core.pct <- core.pct[, colSums(core.pct > 5) > 0]
core.data <- melt(core.pct, id = c("depth", "age"))
ggplot(data = core.data, aes(x = value, y = age)) + geom_path(alpha = 0.5) +
  geom_point() + facet_wrap(~variable, nrow = 1) + scale_y_reverse(expand = c(0,
  0)) + scale_x_continuous(breaks = c(0, 25, 50, 75), expand = c(0, 0)) +
  xlab("Percent Pollen") + ylab(all.downloads[[1]]$chronologies[[1]]$AgeType[1])
```

Andy Lake ((Szeicz *et al.* 1995)) shows changes through time, particularly for *Betula* and *Alnus*, but little *Pinus* pollen.

Pollen data is found in the `counts` component. We want to determine which sample has the first local *Pinus* presence using a cutoff of 5% ((Strong & Hills 2013)). Programmatically we can find which rows in the *Pinus* column have presence over 5% and then find the highest row number since age increases with row number.

```
top.pinus <- function(x) {
  # Convert the core data into proportions by dividing counts by the sum of
  # the row.
  x.pct <- x$counts/rowSums(x$counts)
  # Find the highest row index associated with Pinus presence over 5%
  oldest.row <- max(which(x.pct[, "Pinus"] > 0.05))
  # return a data frame with site name and locations, and then the age and
  # date type associated with the oldest recorded Pinus presence. We preserve
```

```

# date type since some records have ages in radiocarbon years.
data.frame(site = x$metadata$site.data$SiteName, lat = x$metadata$site.data$LatitudeNorth,
  long = x$metadata$site.data$LongitudeWest, age = x$sample.meta$Age[oldest.row],
  date = x$sample.meta$AgeType[oldest.row])
}
# Apply this function to each core (here we use the plyr functions so we can
# return a data.frame instead of a list).
summary.pinus <- ldply(compiled.cores, top.pinus)
# We need to calibrate dates that are recorded in radiocarbon years. In
# most cases we have no idea what the uncertainty was. For this example I
# am simply assuming a 100 year SD for calibration. This is likely too
# conservative.
radio.years <- summary.pinus$date %in% "Radiocarbon years BP"
calibrated <- BchronCalibrate(summary.pinus$age[radio.years], ageSds = rep(100,
  sum(radio.years, na.rm = TRUE)), calCurves = rep("intcal13", sum(radio.years,
  na.rm = TRUE)))
wmean.date <- function(x) sum(x$ageGrid * x$densities/sum(x$densities))
summary.pinus$age[radio.years] <- sapply(calibrated, wmean.date)
# Can be improved by assuming a monotone smooth spline.
regress <- ggplot(summary.pinus, aes(x = lat, y = age)) + geom_point(aes(color = long),
  size = 2) + scale_y_reverse(expand = c(0, 100)) + xlab("Latitude North") +
  ylab("Years Before Present") + geom_smooth(n = 40, method = "loess") + geom_rect(aes(xmin = 59,
  xmax = 60, ymin = 7000, ymax = 10000), color = 2, fill = "blue", alpha = 0.01)
mapped <- ggmap(bc.map) + geom_point(data = summary.pinus, aes(x = long, y = lat,
  color = age), size = 2)
grid.arrange(mapped, regress, nrow = 1)

```

And so we see a clear pattern of migration by *Pinus* in northwestern North America. These results match up broadly with the findings of Strong and Hills ((Strong & Hills 2013)) who suggest that *Pinus* reached a northern extent between 59 and 60°N at approximately 7 - 10kyr as a result of geographic barriers.

Mammal Distributions in the Pleistocene

Graham et al. (Graham *et al.* 1996) look for patterns of change in mammal distributions through the Pleistocene to modern era using fossil assemblages collated from FAUNMAP. The paper uses multiple complex analyses to show in part, that mammal species have responded in a Gleasonian manner to climate change since the late-Pleistocene. Their paper shows some species migrating northward in response to warming climates, others staying relatively stable and some moving southward. Since FAUNMAP has been incorporated into Neotoma we aim to replicate tests of species distributional changes in a straightforward manner to demonstrate the utility of *neotoma* in analysing mammal distributions and change through time.

First we need to obtain all fossil assemblages from Neotoma for vertebrate fauna,

```

# Bounding box is effectively the continental USA, excluding Alaska.
mam.set <- get_dataset(datasettype = "vertebrate fauna", loc = c(-125, 24, -66,
  49.5))
# Calling this many sites can be very time consuming. It takes
# approximately an hour to run fully.
mam.dl <- get_download(dataset = mam.set)

```

So, now we have all the sites, we need to bin them into time periods as in Graham et al. (Graham *et al.* 1996). To do that we first need to build a large table with time and xy coordinates for each site. Time data in *sample.meta* is not the same as for pollen data, where many pollen sites contain an age (often mean

age) and upper and lower bounds. Most mammal sites have younger and older bounds, but no estimates of exact age. In this case we take a short-cut and simply average the younger and older bounds to save the reader from having to examine too much code.

```
library("plyr")
# To be moved into the neotoma package.
source("R/compile_it.R")
compiled.mam <- ldply(mam.dl, .fun = compile_it, .progress = "text")
# We assign time bins to the data. The command findInterval should tell us
# if it is in an interval equivalent to the Modern (0 - 500ybp), Late
# Holocene (500 - 4000ybp), Early-Mid Holocene (4kyr - 10kyr), Late Glacial
# (10kyr - 15kyr), Full Glacial (15kyr - 20kyr) or Late Pleistocene
# (20kyr+).
time.bins <- c(500, 4000, 10000, 15000, 20000)
# This is not the best option, age bounds cross our pre-defined bins,
# however solving this is more complex than this example requires.
mean.age <- rowMeans(compiled.mam[, c("ageold", "ageyoung", "age")], na.rm = TRUE)
interval <- findInterval(mean.age, time.bins)
periods <- c("Modern", "Late Holocene", "Early-Mid Holocene", "Late Glacial",
             "Full Glacial", "Late Pleistocene")
compiled.mam$ageInterval <- periods[interval + 1]
mam.melt <- melt(compiled.mam, measure.vars = 10:(ncol(compiled.mam) - 1), na.rm = TRUE,
                 factorsAsStrings = TRUE)
mam.melt$ageInterval <- factor(mam.melt$ageInterval, levels = periods)
mam.lat <- dcast(data = mam.melt, variable ~ ageInterval, value.var = "lat",
                 fun.aggregate = mean, drop = TRUE)[, c(1, 3, 5, 6)]
# We only want taxa that appear at all time periods:
mam.lat <- mam.lat[rowSums(is.na(mam.lat)) == 0, ]
# Group the samples based on the range & direction (N vs S) of migration.
mam.lat$grouping <- factor(findInterval(mam.lat[, 2] - mam.lat[, 4], c(-11,
-1, 1, 20)), labels = c("Southward", "Stationary", "Northward"))
mam.lat.melt <- melt(mam.lat)
colnames(mam.lat.melt)[2:3] <- c("cluster", "Era")

ggplot(mam.lat.melt, aes(x = Era, y = value)) + geom_path(aes(group = variable,
color = cluster)) + facet_wrap(~cluster) + scale_x_discrete(expand = c(0.1,
0)) + theme(axis.text.x = element_text(angle = 90, hjust = 1))
```

So we can see that at this basic analytic scale species are not uniformly responding to climatic warming following deglaciation. These findings basically echo those of Graham et al. (???) who showed that taxon response is largely individualistic. While we do see the pre-ponderance of migration is northward, a number of taxa show little migratory response and a number show southward migration. In this example we fail to include movement to the west or east, and ignore the issues that may be associated with the complex topography of the mountainous west. Regardless, it is clear that the use of **neotoma** can support research that is reproducible and robust.

Conclusion

The increasing pressure to develop large-scale databases requires the development of tools that can access the data and can leave reproducible analyses so that others can build from and verify results.

Here we present the **neotoma** package for R (R Core Team 2014) and use examples from the literature to show its utility. **neotoma** joins a number of other existing packages that are designed either to exploit existing

paleoecological datasets (Group 2014) or to manipulate paleoecological data (Simpson 2007; Juggins 2013; Simpson & Oksanen 2013). The **neotoma** package itself is available either from the CRAN repository, or from GitHub where ongoing development continues with help from the public.

The use of the Neotoma database continues to expand, and here we provide researchers with the tools to move analytics to an open framework using R (R Core Team 2014) so that methods can be more fully visible.

1. Behrensmeier, A.K. & Miller, J.H. *et al.* (2012). Building links between ecology and paleontology using taphonomic studies of recent vertebrate communities. In: *Paleontology in ecology and conservation*. Springer, pp. 69–91.
2. Blaauw, M., Christen, J.A. & others *et al.* (2011). Flexible paleoclimate age-depth models using an autoregressive gamma process. *Bayesian Analysis*, 6, 457–474.
3. Blarquez, O., Carcaillet, C., Frejaville, T. & Bergeron, Y. *et al.* (2014). Disentangling the trajectories of alpha, beta and gamma plant diversity of north american boreal ecoregions since 15,500 years. *Frontiers In Paleoecology*, 2, 6.
4. Blois, J.L., Williams, J.W., Fitzpatrick, M.C., Ferrier, S., Veloz, S.D. & He, F. *et al.* (2013). Modeling the climatic drivers of spatial patterns in vegetation composition since the last glacial maximum. *Ecography*, 36, 460–473.
5. Boettiger, C., Lang, D. & Wainwright, P. *et al.* (2012). rfishbase: exploring, manipulating and visualizing fishBase data from r. *Journal of fish biology*, 81, 2030–2039.
6. Brewer, S., Jackson, S.T. & Williams, J.W. *et al.* (2012). Paleoecoinformatics: applying geohistorical data to ecological questions. *Trends in ecology & evolution*, 27, 104–112.
7. Chamberlain, S.A. & Szöcs, E. *et al.* (2013). taxize: taxonomic search and retrieval in r. *F1000Research*, 2.
8. Dietl, G.P. & Flessa, K.W. *et al.* (2011). Conservation paleobiology: putting the dead to work. *Trends in Ecology & Evolution*, 26, 30–37.
9. Fritz, S.A., Schnitzler, J., Eronen, J.T., Hof, C., Böhning-Gaese, K. & Graham, C.H. *et al.* (2013). Diversity in time and space: wanted dead and alive. *Trends in ecology & evolution*, 28, 509–516.
10. Gavin, D.G., Oswald, W.W., Wahl, E.R. & Williams, J.W. *et al.* (2003). A statistical approach to evaluating distance metrics and analog assignments for pollen records. *Quaternary Research*, 60, 356–367.
11. Gill, J.L., McLauchlan, K.K., Skibbe, A.M., Goring, S., Zirbel, C.R. & Williams, J.W. *et al.* (2013). Linking abundances of the dung fungus *sporormiella* to the density of bison: implications for assessing grazing by megaherbivores in palaeorecords. *Journal of Ecology*, 101, 1125–1136.
12. Goring, S., Lacourse, T., Pellatt, M.G. & Mathewes, R.W. *et al.* (2013). Pollen assemblage richness does not reflect regional plant species richness: a cautionary tale. *Journal of Ecology*, 101, 1137–1145.
13. Goring, S., Williams, J., Blois, J., Jackson, S., Paciorek, C. & Booth, R. *et al.* (2012). Deposition times in the northeastern united states during the holocene: establishing valid priors for bayesian age models. *Quaternary Science Reviews*, 48, 54–60.
14. Graham, R.W., Lundelius Jr, E.L., Graham, M.A., Schroeder, E.K., Toomey III, R.S. & Anderson, E. *et al.* (1996). Spatial response of mammals to late quaternary environmental fluctuations. *Science*, 272, 1601–1606.
15. Group, G.P.W. *et al.* (2014). *paleofire: paleofire: an r package to analyse sedimentary charcoal records from the global charcoal database to reconstruct past biomass burning*.
16. Juggins, S. *et al.* (2013). *rioja: Analysis of quaternary science data*.
17. MacDonald, G. & Cwynar, L.C. *et al.* (1991). Post-glacial population growth rates of *pinus contorta* ssp. *latifolia* in western canada. *The Journal of Ecology*, 417–429.
18. Parnell, A. *et al.* (2014). *Bchron: Radiocarbon dating, age-depth modelling, relative sea level rate estimation, and non-parametric phase modelling*.

- 19.R Core Team *et al.* (2014). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing, Vienna, Austria.
- 20.Simpson, G.L. *et al.* (2007). Analogue methods in palaeoecology: Using the analogue package. *Journal of Statistical Software*, 22, 1–29.
- 21.Simpson, G.L. & Oksanen, J. *et al.* (2013). *analogue: Analogue and weighted averaging methods for palaeoecology*.
- 22.Strong, W.L. & Hills, L.V. *et al.* (2013). Holocene migration of lodgepole pine (*pinus contorta* var. *latifolia*) in southern yukon, canada. *The Holocene*, 23, 1340–1349.
- 23.Szeicz, J.M., MacDonald, G.M. & Duk-Rodkin, A.*et al.* (1995). Late quaternary vegetation history of the central mackenzie mountains, northwest territories, canada. *Palaeogeography, Palaeoclimatology, Palaeoecology*, 113, 351–371.
- 24.Uhen, M.D., Barnosky, A.D., Bills, B., Blois, J., Carrano, M.T. & Carrasco, M.A.*et al.* (2013). From card catalogs to computers: databases in vertebrate paleontology. *Journal of Vertebrate Paleontology*, 33, 13–28.